

1 Word embeddings

Open the ipython notebook `word_embedding_intro_release.ipynb`. Include in your writeup

- a frequency histogram, and a histogram of the co-occurrences of the words in the dataset
- the list of the 10 most frequent words, and the 10 least frequent words
- the list of the 10 most cooccurring word pairs
- the 10 closest words to the city where you were born. (If you were born in a town which is not in the vocabulary, pick the closest city which *is* in the vocabulary.
- the 10 closest words to an object that is close to you right now.
- your 2-D PCA word embeddings, and some comments on any interesting geometric structure you may see

2 Duality theory and SVMs

- Remember the soft-margin SVM problem from forever ago?

$$\begin{aligned} \underset{\theta, s}{\text{minimize}} \quad & \frac{1}{2} \|\theta\|_2^2 + C \sum_{i=1}^m \max\{s_i, 0\} \\ \text{subject to} \quad & y_i x_i^T \theta + s_i = 1, \quad i = 1, \dots, m \end{aligned} \tag{1}$$

We're going to use this method to classify our parts of speech.

- **Subgradients and subdifferentials.** The hinge loss function

$$g(s) = \sum_{i=1}^m \max\{s_i, 0\}$$

is unfortunately not differentiable whenever there exists one $s_i = 0$. But, it is convex. So, we can still take descent steps with respect to the objective by looking for *subgradients*. Specifically, for the scalar function

$$g_i(s) = \max\{s_i, 0\},$$

we describe the *subdifferential* of g_i at s as

$$\partial g_i(s) = \mathbf{conv} \left\{ \lim_{\epsilon \rightarrow 0^+} g'_i(s + \epsilon), \lim_{\epsilon \rightarrow 0^-} g'_i(s + \epsilon) \right\} = \begin{cases} \{1\} & s_i > 0 \\ \{0\} & s_i < 0 \\ [0, 1] & s_i = 0. \end{cases}$$

Since subdifferentials are linear (not obvious, but can be proved), we can summarize the subdifferential of the hinge function via

$$\partial(C \cdot g(s)) = C \sum_{i=1}^m \partial g_i(s)$$

where we describe the transformation of sets as

$$CS = \{Cx : x \in \mathcal{S}\} \text{ and } \mathcal{S}_1 + \mathcal{S}_2 = \{x + y : x \in \mathcal{S}_1, y \in \mathcal{S}_2\}.$$

- The subgradient method for minimizing $f(x)$ for some convex but not differentiable function f goes as

$$x^{(t+1)} = x^{(t)} - \alpha^{(t)} g^{(t)}, \quad g^{(t)} \in \partial f(x^{(t)}).$$

Here, $g^{(t)}$ is any element from $\partial f(x^{(t)})$. When you have more than one choice, you can pick any element, but you should design a rule that is agnostic to your knowledge of where the true solution is. When $\alpha^{(t)}$ is diminishing but not summable, e.g.

$$\alpha^{(t+1)} < \alpha^{(t)}, \quad \lim_{t \rightarrow +\infty} \sum_{i=1}^t \alpha^{(i)} = \infty$$

then the subgradient method is known to converge. In particular, picking $\alpha^{(t)} = 1/(Lt)$, the method converges at a rate of $f(x^{(t)}) - f^* = O(1/\sqrt{t})$.

- Q1** Show that for a *constant* step size, subgradient descent doesn't converge. Show this by arguing that minimizing the function

$$\underset{x}{\text{minimize}} \quad \frac{1}{2}(x - c)^2 + \max\{x, 0\}$$

cannot reach its optimum for *all* choices of c , using a step size that is agnostic to c .

- **Projections.** Consider the projection problem

$$\begin{aligned} \underset{x}{\text{minimize}} \quad & \frac{1}{2} \|x - \hat{x}\|_2^2 \\ \text{subject to} \quad & Ax = b \end{aligned} \tag{2}$$

- Q2** Find the Lagrange dual of (2) and use the solution of the dual to show that, assuming that AA^T is invertible, then the solution to (2) is

$$x = A^T(AA^T)^{-1}(b - A\hat{x}) + \hat{x}.$$

- **Projected subgradient descent.** For a problem of form

$$\begin{aligned} \underset{x}{\text{minimize}} \quad & f(x) \\ \text{subject to} \quad & Ax = b \end{aligned}$$

the projected subgradient descent method runs the iterative scheme (from any initial point)

$$x^{(t+1)} = \text{proj}_{\mathcal{H}}(x^{(t)} - \alpha g^{(t)})$$

where $g^{(t)}$ is a subgradient of f at $x^{(t)}$ and $\mathcal{H} = \{x : Ax = b\}$ and the projection of \hat{x} on \mathcal{H} is the solution to problem (2).

- Q3** Use the pieces you have so far to propose an iteration scheme to solve (1) via projected subgradient descent. That is, given some $\theta^{(t)}$ and $s^{(t)}$, write out the equations used to compute $\theta^{(t+1)}$ and $s^{(t+1)}$, as explicitly as possible. (You will implement this code in the next section, so do it carefully here.)

- An equivalent formulation of (1) is the following

$$\begin{aligned} \underset{\theta, b, s}{\text{minimize}} \quad & \frac{1}{2} \|\theta\|_2^2 + C \sum_{i=1}^m s_i \\ \text{subject to} \quad & y_i x_i^T \theta + s_i \geq 1, \quad i = 1, \dots, m \\ & s_i \geq 0. \end{aligned} \tag{3}$$

- Q4** Write down the Lagrangian saddle point problem of (3) and minimize it with respect to the primal variables to derive the following dual problem

$$\begin{aligned} \underset{u}{\text{maximize}} \quad & -\frac{1}{2} \sum_{i=1}^m (y_i x_i^T u)^2 + \sum_{i=1}^m u_i \\ \text{subject to} \quad & 0 \leq u \leq 1. \end{aligned} \tag{4}$$

Q5 Show that if θ^* , s^* optimize (3) and u^* optimizes (4), then

1. $s_i^* > 1$ implies the i th training sample is not classified correctly,
2. $y_i x_i^T \theta > 1$ implies $u_i = 0$,
3. $y_i x_i^T \theta < 1$ implies $u_i = C$,
4. $0 < u_i < C$ implies $y_i x_i^T \theta = 1$.

Q6 Given u^* , propose a predictor, e.g. some function where $f(u^*, x) = x^T \theta^*$.

Q7 Propose how you would solve this using projected gradient descent, and how you would form a predictor. Remember that later, you will code this up, so include any detail here you would need to write this code.

- **Kernel support vector machines** One way to think about kernel functions is to replace the linear prediction

$$y = \text{sign}(x_i^T \theta)$$

with a lifted version of x_i , as $\phi(x_i)$, e.g.

$$y = \text{sign}(\phi(x_i)^T \theta).$$

After finding the dual, this gives us an objective function of

$$g(u) = -\frac{1}{2} \sum_{i=1}^m (y_i \phi(x_i)^T u)^2 + \sum_{i=1}^m u_i = -\frac{1}{2} u^T K u + u^T \mathbf{1}$$

where $K_{ij} = y_i y_j \phi(x_i)^T \phi(x_j)$ defines the Kernel matrix. The Kernel trick is simply, instead of holding onto the representations $\phi(x)$, to only hold onto these inner product values. For example, the radial basis function (RBF) kernel computes

$$K_{ij} = K(x_i, x_j) = \exp\left(-\frac{\|x_i - x_j\|_2^2}{2\sigma^2}\right), \quad \hat{K}_{ij} = K_{ij} y_i y_j$$

which can be precomputed.

$$\begin{aligned} & \underset{u}{\text{maximize}} && -\frac{1}{2} \sum_{i=1}^m u^T \hat{K} u^2 + \sum_{i=1}^m u_i \\ & \text{subject to} && 0 \leq u \leq 1. \end{aligned} \tag{5}$$

Q8 Propose how, given u^* , you would offer a prediction on a new datasample, x . Remember that you do not have access to ϕ , but you do have access to K .

Q9 Propose how you would solve this using projected gradient descent. Remember that later, you will code this up, so include any detail here you would need to write this code.

3 Multiclass classification

Now open `svm_POS_release.ipynb` and go through the notebook. Report in your writeup

- loss, train and test misclassification plots for noun-vs-all, for primal projected subgradient method
- comparison of loss functions for primal and dual linear SVM
- evidence of cross validation to determine best value of C in the primal linear SVM and C and σ in Kernel SVM
- confusion matrix of final answer on train, validation, and test set.

4 Discussion

Conclude your report with some observations and thoughts on

- the computational complexity of each method
- the effectiveness of each method
- any tricks used to deal with imbalanced data, or with overfitting.

Would you recommend using word embedding for POS classification, using the method we proposed here? Using a different method? Not at all?